

# Red Belly Blockchain

Vincent Gramoli

(Special thanks to Tyler Crain, Mikel Larrea, Chris Natoli, Michel Raynal, Guillaume Vizier)



THE UNIVERSITY OF  
SYDNEY





# Sydney





# Red Belly Snake

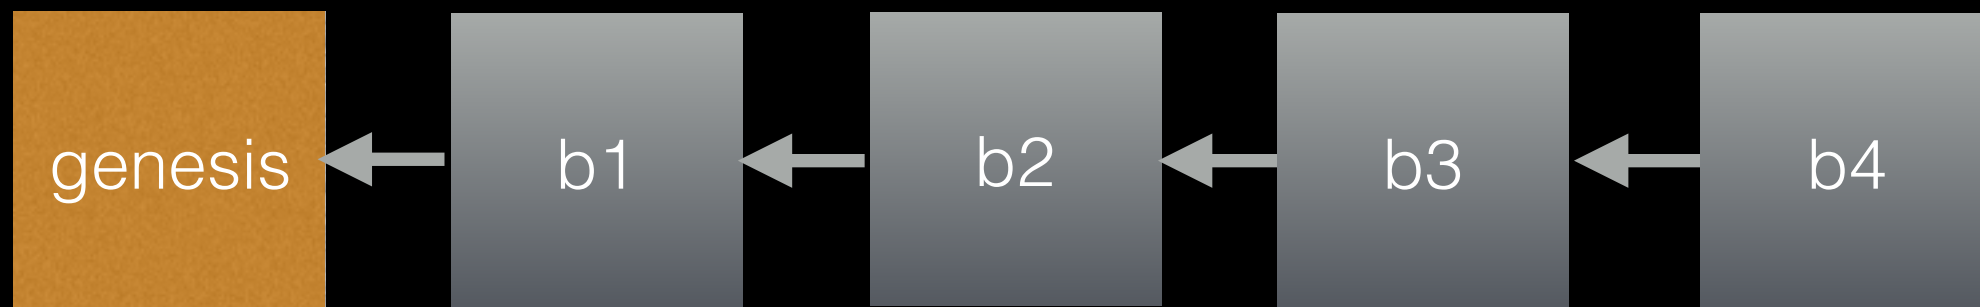


# Roadmap

- Context: Blockchain
- Problem: Balance Attack
- Solution: Unforkable Blockchain
  - Democratic BFT consensus
  - Red Belly Blockchain

What is a blockchain?

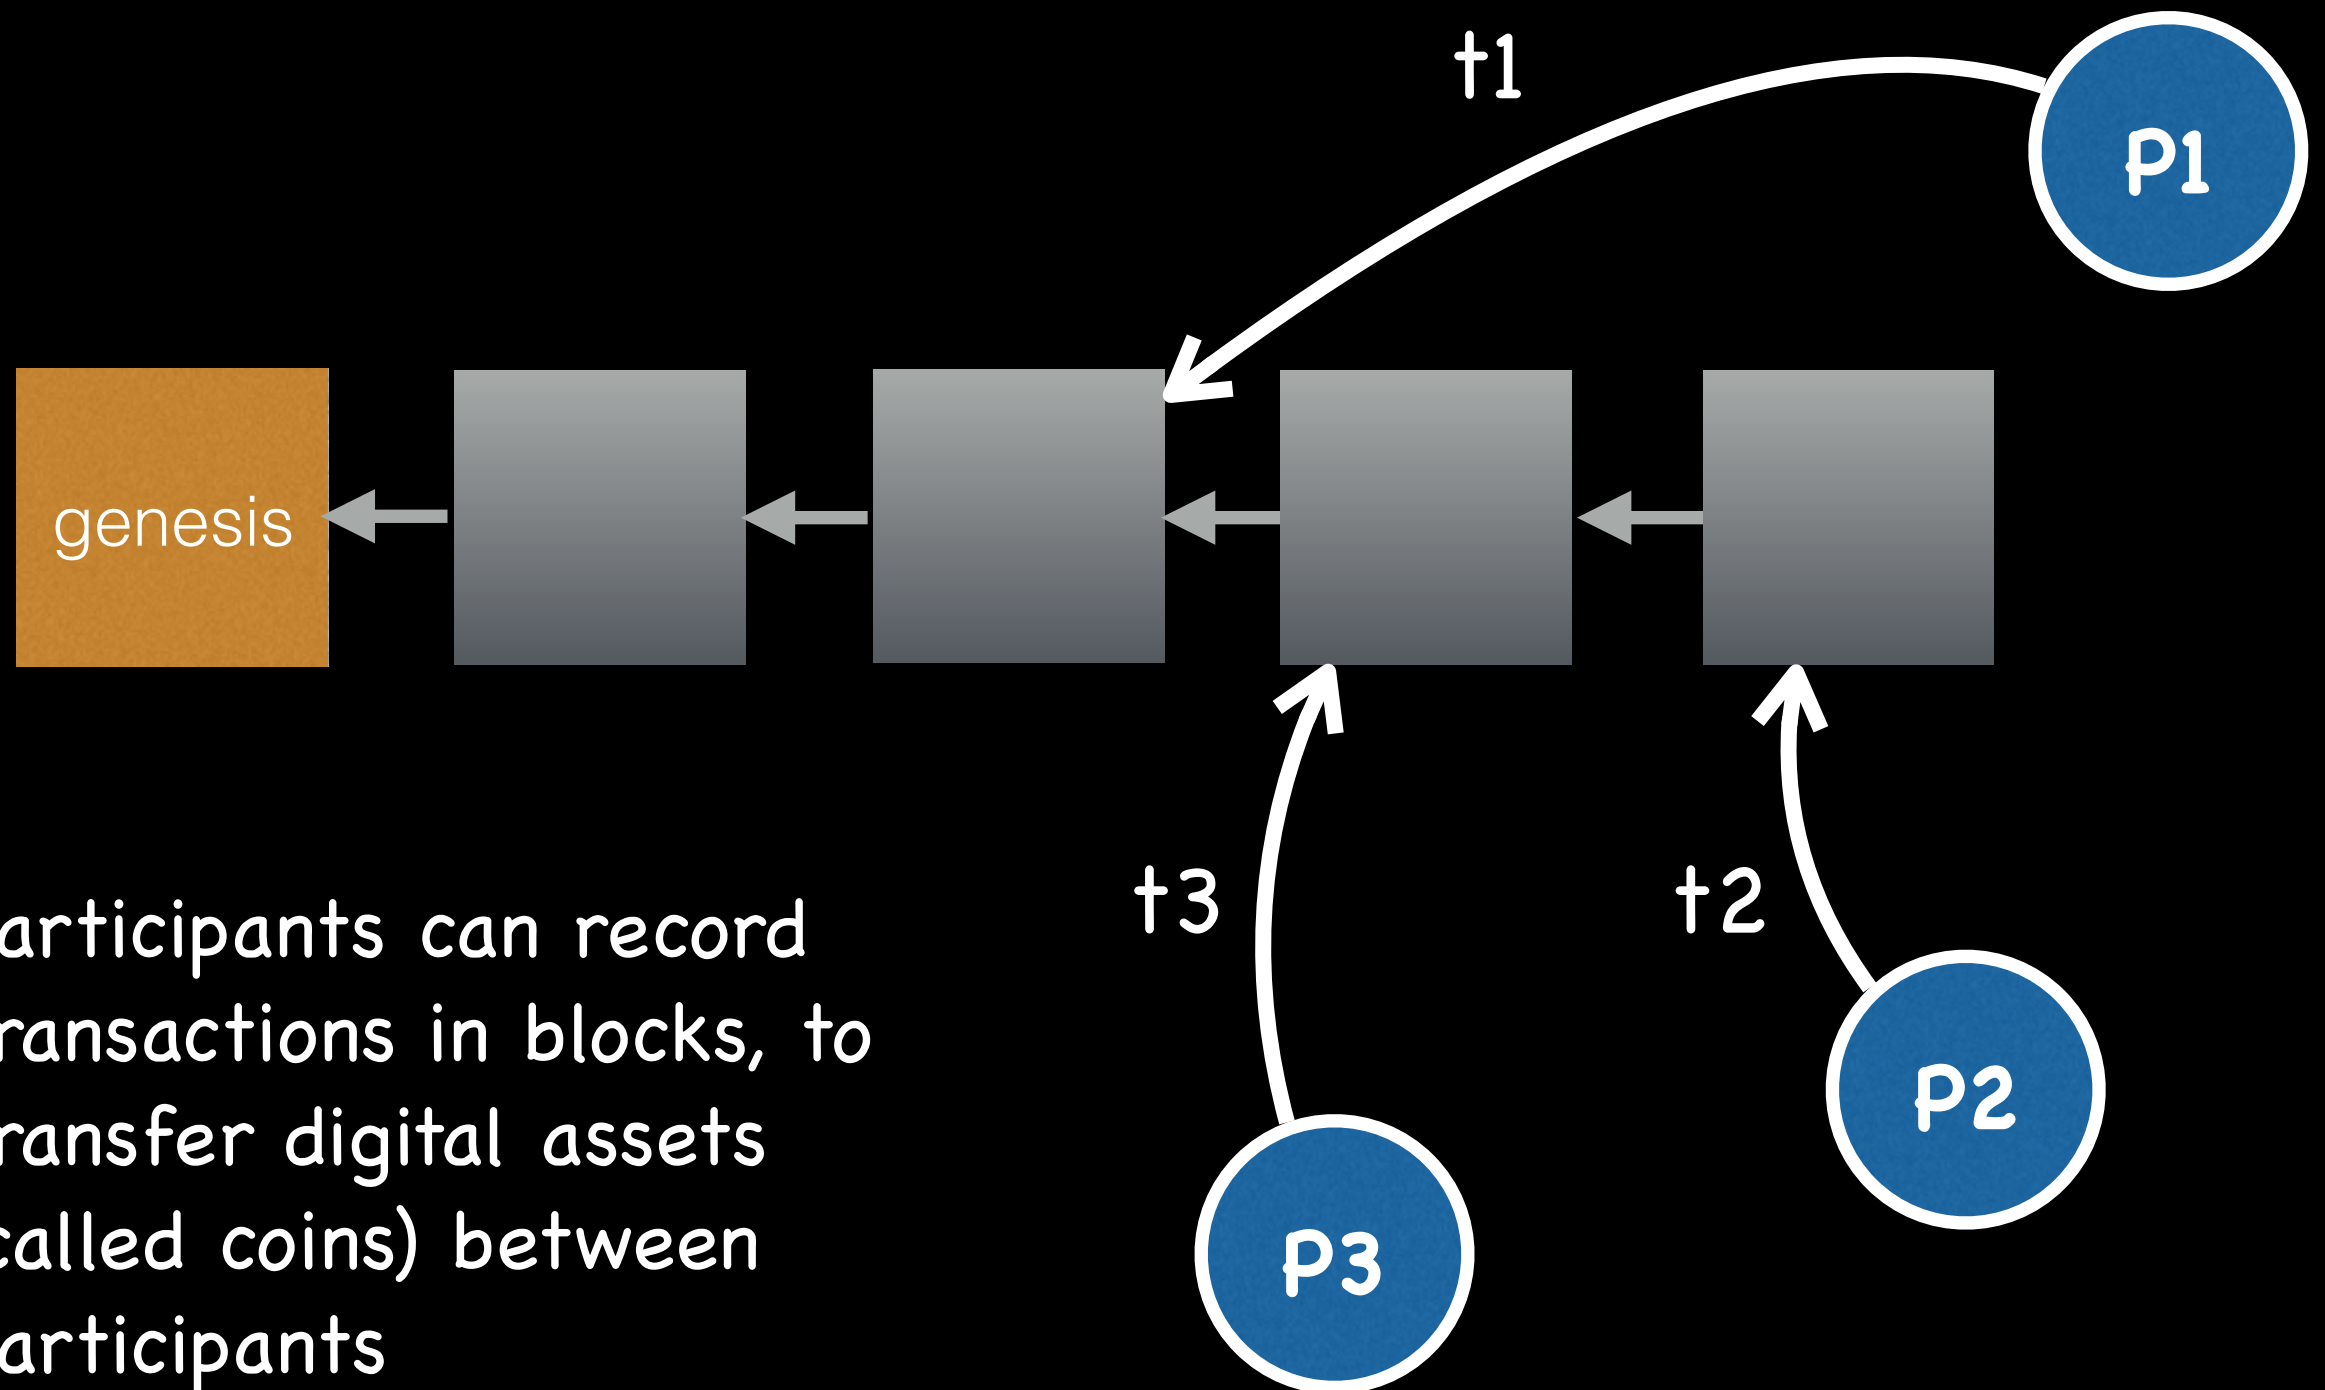
# Blockchain [Nak'08]



- Chain of blocks
- Starts at a special block called the genesis block
- Each block contains the hash of its predecessor
- To create a block one must solve a crypto-puzzle




# Blockchain (con't)




Participants can record transactions in blocks, to transfer digital assets (called coins) between participants

# Blockchain (con't)

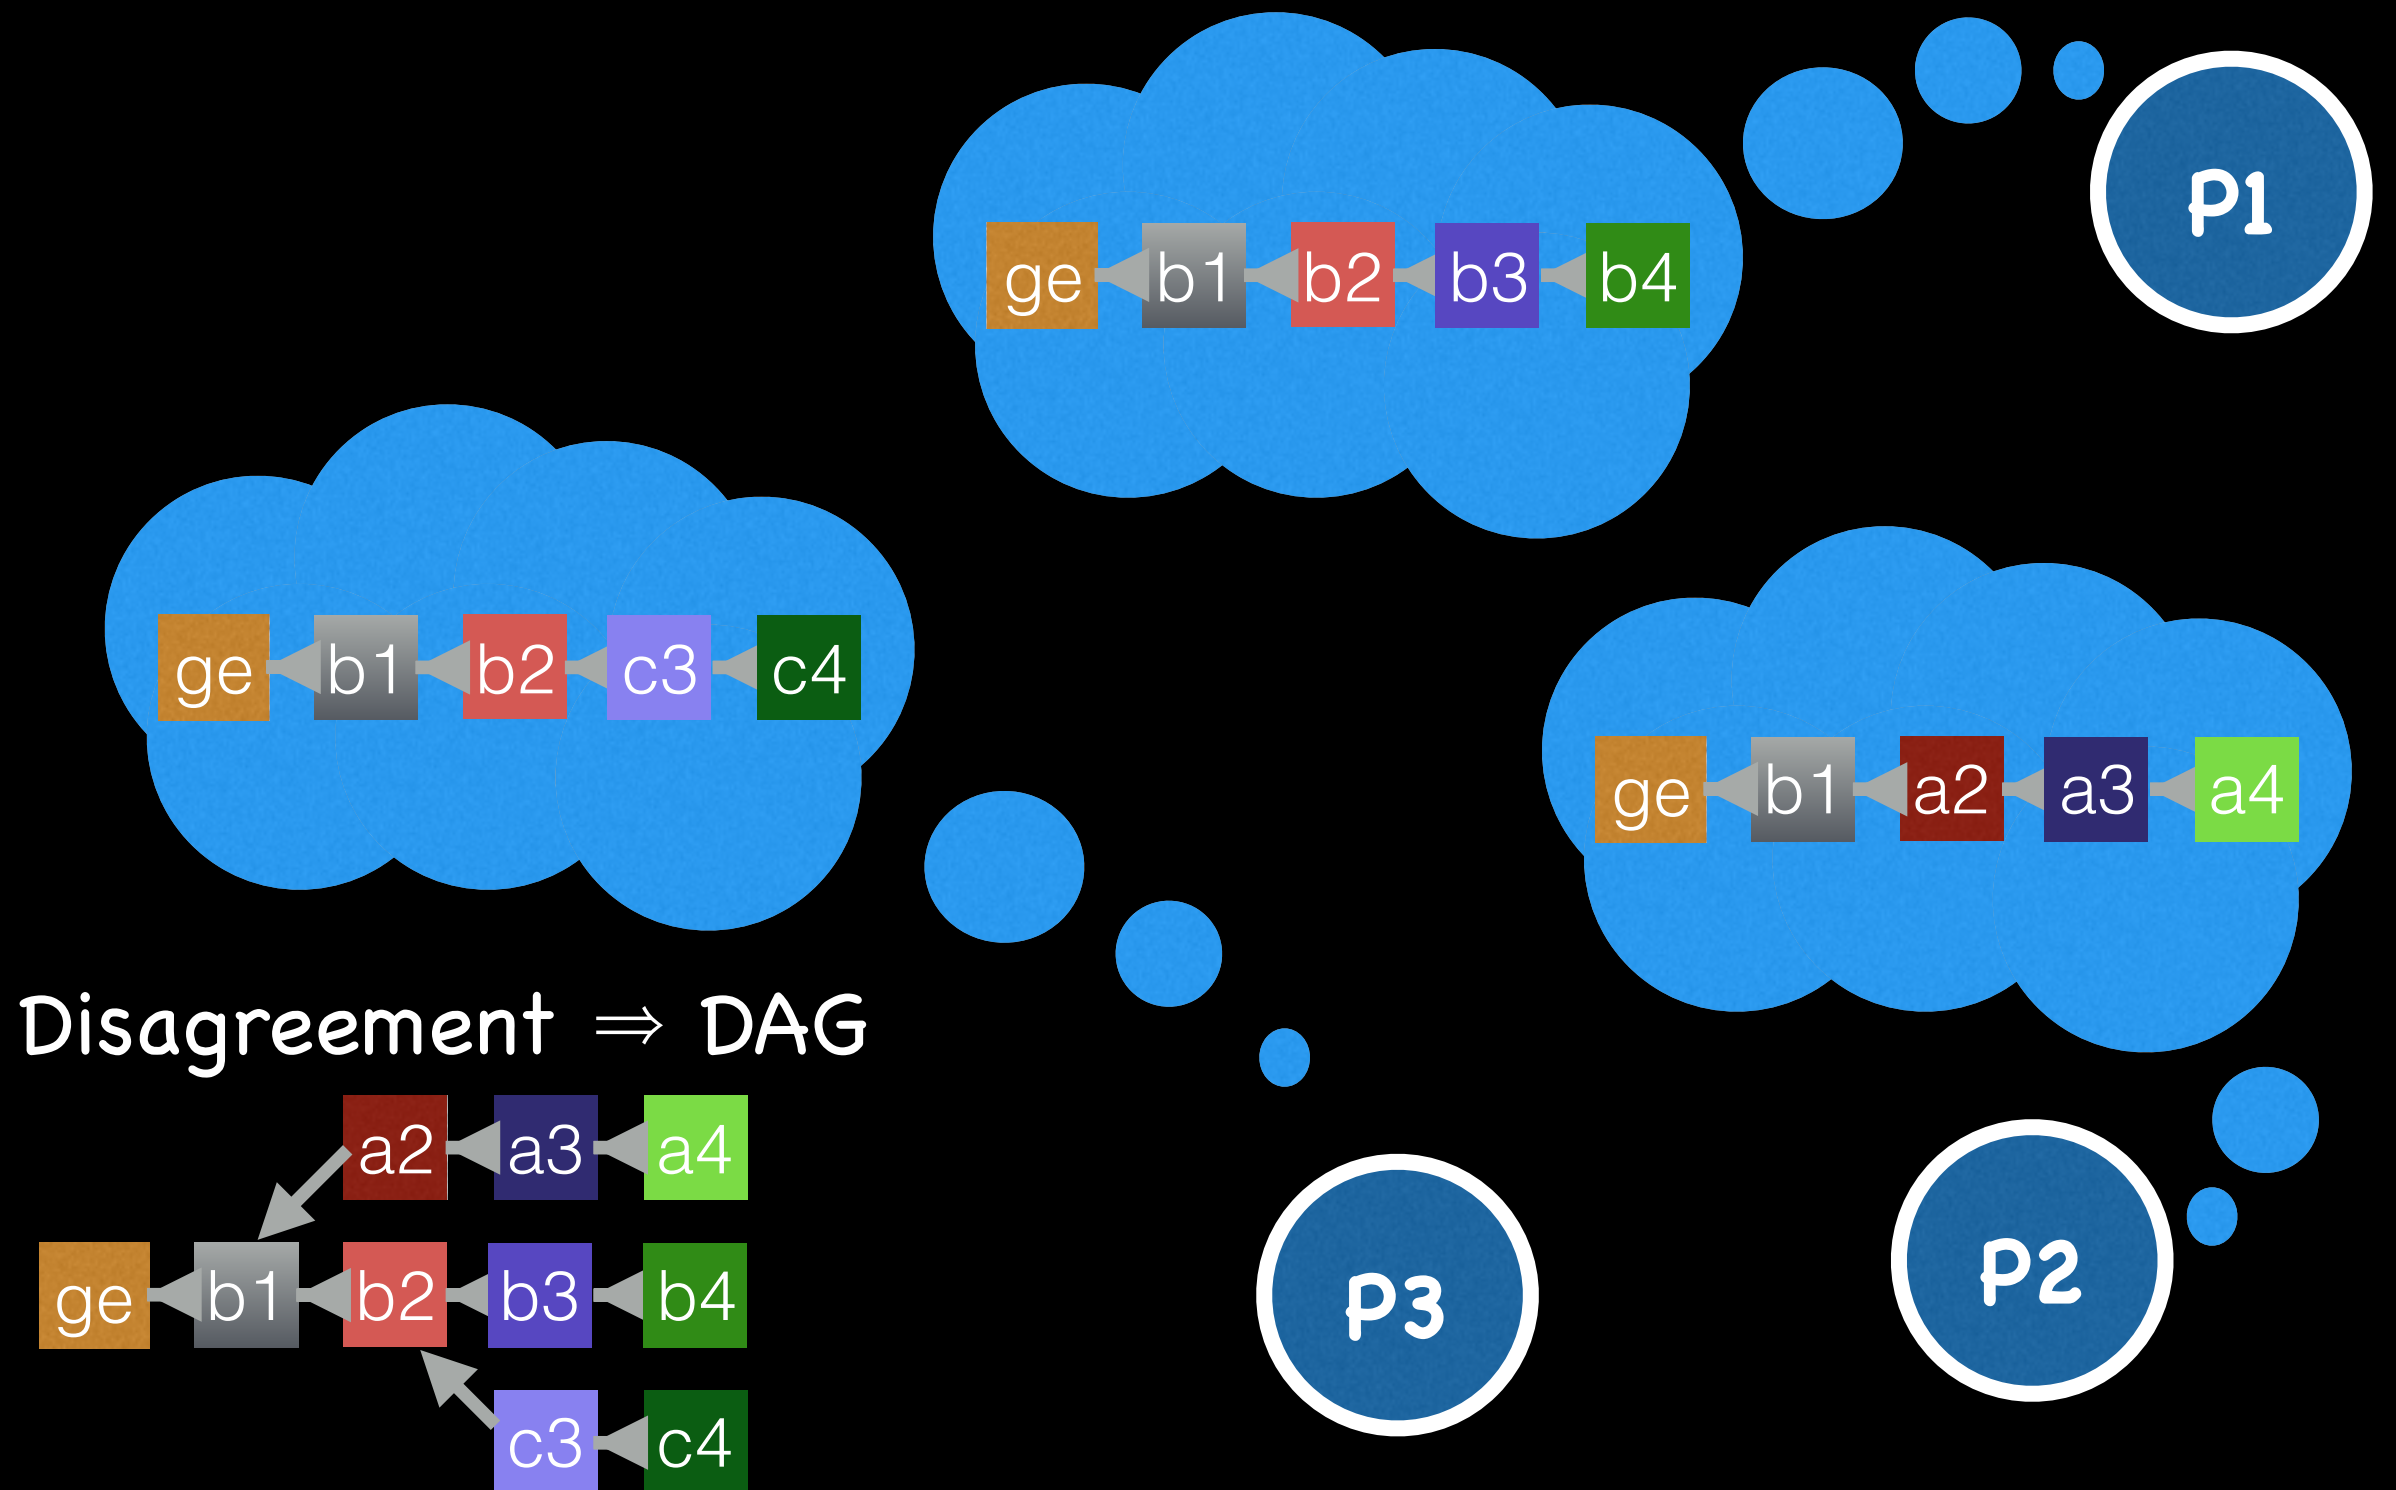





# Blockchain (con't)



# Blockchain (con't)




# Blockchain (con't)






# Blockchain (con't)

$\sigma_A(t_A)$ : Alice gives all her coins to **Bob**



$\sigma_A(t_{A'})$ : Alice gives all her coins to **Carole**

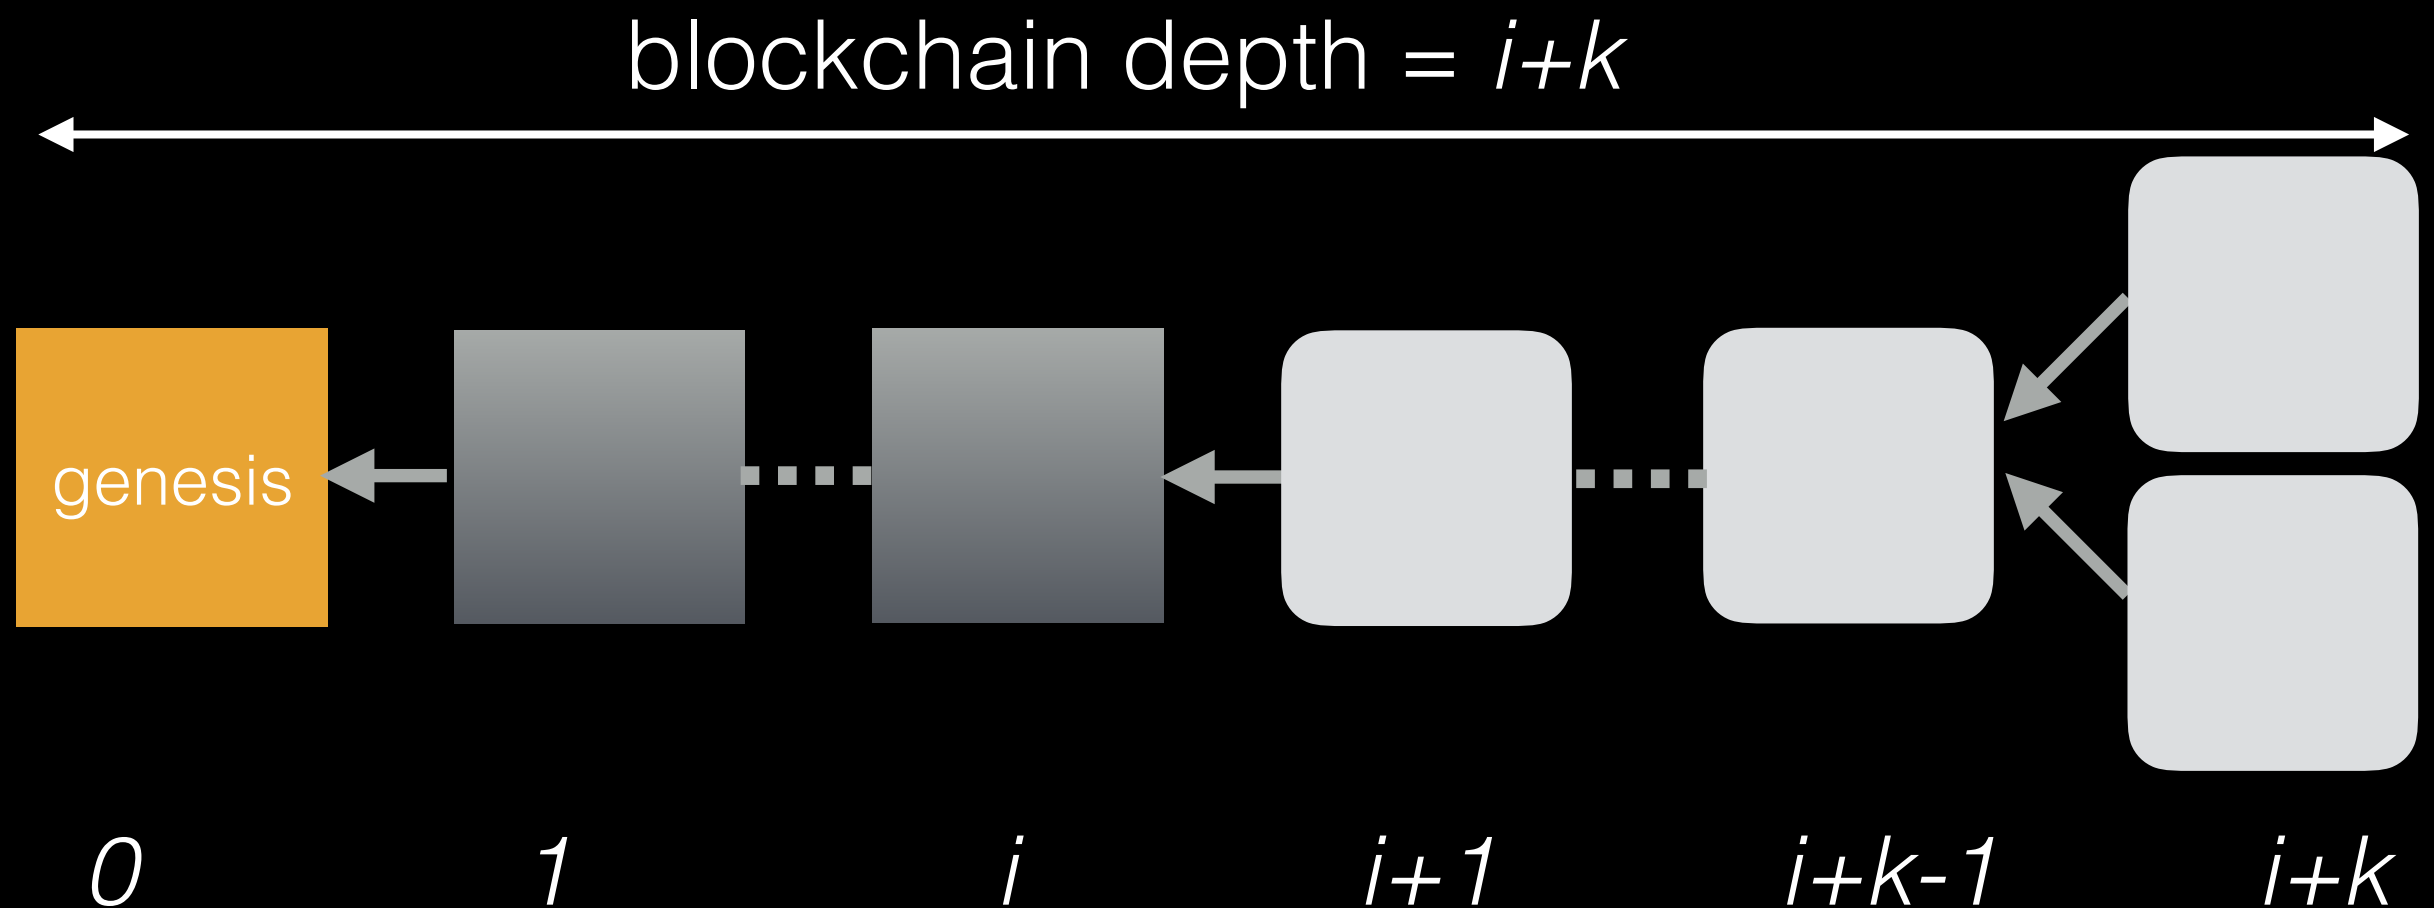
# Blockchain (con't)



One branch is selected

based on its length, the weight of its subtrees, its content...

# Blockchain (con't)



 genesis block     decided block     undecided block


We say that a transaction commits when it is in a decided block [NCA'16]



What's dangerous  
about it?


# The Balance Attack [DSN'17]

Consider a communication graph  $G=\langle V,E\rangle$



# The Balance Attack (con't)

Consider a communication graph  $G=\langle V,E\rangle$




Let  $G_1$  and  $G_2$  be two subgraphs of  $G$  with the same mining power separated by  $E_3$  such that  $E = E_1 \cup E_2 \cup E_3$



# The Balance Attack (con't)


Let an attacker issue  $t_A$  in  $G1$  and delay links  $E3$



Then the two subgraphs end up with branches of similar lengths or weights  $W1$  and  $W2$

# The Balance Attack (con't)



The attacker has only to mine  $|W1-W2|+1$  blocks in G2



The weight of a branch in G2 exceeds G1's branches

# The Balance Attack (con't)

When the attacker stops delaying the messages




$t_A$  is discarded, allowing double spending with  $t_A'$



# Unforkable Blockchain

# Unforkable blockchain



We say that a transaction commits when it is in a decided block [NCA'16]

# Model

- **Distributed system:**  $n$  processes  $\Leftarrow$  but additional processes can issue transactions and read the blockchain
- **Partially synchronous:** the upper-bounds on the delay of messages and computation is not known  $\Leftarrow$  Internet can be congested
- **Byzantine failures:** there can only be  $t < n/3$  arbitrary failures, all other processes are correct  $\Leftarrow$  Attackers have incentives to try stealing
- **Point-to-point reliable channel:** a message sent by a correct process is eventually received by a correct process  $\Leftarrow$  Public key encryptions can implement it

# Byzantine Consensus?

Each correct process invokes  $\text{propose}(v)$  with its value  $v$  and decides the returned value such that:

1

Agreement: no two correct processes decide differently

2

Termination: every correct process decides

3

Validity: the decided value is proposed by a correct process



# Byzantine Consensus?

Each correct process invokes  $\text{propose}(v)$  with its value  $v$  and decides the returned value such that:

1

Agreement: no two correct processes decide differently

2

Termination: every correct process decides

3

Validity: the decided value is proposed by a correct process

# Blockchain Consensus [AlgoTel'17]

Provided an application-specific `valid()` predicate, each correct process invokes `propose(v)` and decides the returned value such that:

1

Agreement: no two correct processes decide differently

2

Termination: every correct process decides

3

Validity: the decided value satisfies the predicate `valid()`

# Democratic BFT

# Binary Byzantine Consensus

Each correct process invokes  $\text{propose}(v)$  with its value  $v \in \{0,1\}$  and decides the returned value such that:

1

Agreement: no two correct processes decide differently

2

Termination: every correct process decides

3

Validity: if all correct processes propose the same value, then no other value can be decided



# Safe DBFT [CGLR'17]

est: local current estimate of the value to decide

r: local round number, initially 0

bin-values[]: array of binary values for all rounds

b, auxiliary binary value

values<sub>i</sub>, auxiliary set of values

The algorithm uses two message types

EST[r]() used at round r to BV-bcast est.

AUX[r]() used at round r to broadcast bin-values[r]

# Safe DBFT (con't)

## BV-broadcast [JACM'15]

- BV-obligation: if  $t+1$  correct BV-bcast  $v$ , then  $v$  is eventually added to the set bin-values of all correct processes
- BV-justification: if  $p_i$  is correct and  $v$  in bin-values $_i$ , then  $v$  was BV-bcast by a correct process
- BV-uniformity: if  $v$  is added to bin-values $_i$  of correct  $p_i$ , then eventually  $v$  will be in bin-values $_j$  for all correct  $p_j$
- BV-termination: eventually bin-values of correct  $p_i$  is not empty

# Safe DBFT (con't)

$est \leftarrow v$

$r \leftarrow 0$

**while** (**true**) {

$r \leftarrow r+1$

    BV-bcast EST[r](est)

**wait until** (bin-values[r]  $\neq \emptyset$ )

    bcast AUX[r](bin-values[r])

**wait until** (messages AUX[r](bval<sub>p<sub>1</sub></sub>), ..., AUX[r](bval<sub>p<sub>(n-t)</sub></sub>) received

        from (n-t) distinct processes p(x),  $1 \leq x \leq n-t$ , and their content is such that:

$\exists$  a non-empty set values<sub>i</sub> such that (i) values<sub>i</sub> in bin-values[r]<sub>i</sub> and (ii) values<sub>i</sub> =  $\cup b\_val_x$

$b \leftarrow r \bmod 2$

**if** (values<sub>i</sub> = {v}) {

$est \leftarrow v$

**if** (v=b) **then** decide(v) if not done yet

    } **else** {

$est \leftarrow b$

    }

}

# Safe DBFT (con't)

$est \leftarrow v$

$r \leftarrow 0$

**while** (**true**) {

$r \leftarrow r+1$

BV-bcast EST[r](est)

**wait until** (bin-values[r]  $\neq \emptyset$ )

bcast AUX[r](bin-values[r])

**wait until** (messages AUX[r](bval<sub>p<sub>1</sub></sub>), ..., AUX[r](bval<sub>p<sub>(n-t)</sub></sub>) received

from (n-t) distinct processes p(x),  $1 \leq x \leq n-t$ , and their content is such that:

$\exists$  a non-empty set values<sub>i</sub> such that (i) values<sub>i</sub> in bin-values[r]<sub>i</sub> and (ii) values<sub>i</sub> =  $\cup b\_val_x$

$b \leftarrow r \bmod 2$

**if** (values<sub>i</sub> = {v}) {

$est \leftarrow v$

**if** (v=b) **then** decide(v) **if not done yet**

} **else** {

$est \leftarrow b$

}

}

Filter out values  
proposed only by  
Byzantine processes

# Safe DBFT (con't)

Lemma 1: If at the beginning of a round, all correct processes have the same estimate, they never change their estimate thereafter.

Proof sketch. Assume all correct have the same estimate  $v$  at round  $r$ . BV-Obligation and BV-Justification  $\Rightarrow \text{bin\_values}[r] = \{v\}$ . Hence, at every correct we have  $\text{values} = \{v\}$ , so that  $\text{est}$  becomes  $v$ .



# Safe DBFT (con't)

Lemma 2: Let  $p_i$  and  $p_j$  be correct. If their values are singletons, they are the same.

Proof sketch. If a correct process has values =  $\{v\}$  then it received  $AUX[r]\{v\}$  from  $n-t$  distinct processes and  $t+1$  correct. For two correct processes to have  $w$  and  $v$  as this singleton, it would mean they received these distinct values from  $n-t$  processes each. As  $(n-t)+(t+1) > n$ , one correct process must have sent the same value to both, and we have  $v=w$ .

# Safe DBFT (con't)

Theorem [Agreement]: No two correct processes decide different values.

Proof sketch. Let  $r$  be the 1st round where a correct process decides, hence  $\text{values}[r] = \{v=r \bmod 2\}$ . If another correct process decides  $w$  in the same round, then  $v=w$  by Lemma 2. Let  $p_j$  be a correct that does not decide in round  $r$ . By Lemma 2,  $\text{values}_j \neq \{w\}$  so  $\text{values}_j = \{0,1\}$ . Thus  $\text{est}_j = v$ . All correct have thus the same estimate in round  $r+1$ , and stick to it (Lemma 1).

# Safe DBFT (con't)

BC, binary consensus instance

```
mv-propose() { // similar to [PODC'94]
  RB-bcast VAL(v) // reliable broadcast [Bra'87]
  repeat if ( $\exists k : \text{proposals}[k] \neq \perp \wedge \text{BC}[k].\text{binpropose}() \text{ not invoked}$ )
    BC[k].binpropose(1)
  until( $\exists l : \text{bin-decisions}[l] = 1$ )
  for each k such that BC[k].binpropose() not invoked:
    BC[k].binpropose(0)
  wait until ( $\wedge 1 \leq x \leq n \text{ bin-decisions}[x] \neq \perp$ )
  j = min{x : bin-decisions[j] = 1}
  wait until proposals[j]  $\neq \perp$ 
  decide( $\bigcup_{\forall j: \text{bin-decisions}[j]=1} \text{proposals}[j]$ )
}

when val(v) is RB-delivered from pj do // reliable broadcast delivery
  if valid(v) then
    proposals[j]  $\leftarrow v$ ;
    BV-deliver b-val[1](1) to BC[j]

when BC[k].binpropose() returns b do bin-decisions[k]  $\leftarrow b$ 
```

# Safe DBFT (con't)

BC, binary consensus instance

```
mv-propose() { // similar to [PODC'94]
  RB-bcast VAL(v) // reliable broadcast [Bra'87]
  repeat if ( $\exists k : \text{proposals}[k] \neq \perp \wedge \text{BC}[k].\text{binpropose}() \text{ not invoked}$ )
    BC[k].binpropose(1)
  until( $\exists l : \text{bin-decisions}[l] = 1$ )
  for each k such that BC[k].binpropose() not invoked:
    BC[k].binpropose(0)
  wait until ( $\wedge 1 \leq x \leq n \text{ bin-decisions}[x] \neq \perp$ )
  j = min{x : bin-decisions[j] = 1}
  wait until proposals[j]  $\neq \perp$ 
  decide( $\bigcup_{\forall j: \text{bin-decisions}[j]=1} \text{proposals}[j]$ )
}
```

Spawn multiple  
binary cons.  
instances

Use binary  
decisions as a  
bitmask

```
when val(v) is RB-delivered from pj do // reliable broadcast delivery
  if valid(v) then
    proposals[j]  $\leftarrow v$ ;
    BV-deliver b-val[1](1) to BC[j]
```


```
when BC[k].binpropose() returns b do bin-decisions[k]  $\leftarrow b$ 
```

# Red Belly Blockchain



# The Red Belly Blockchain


All nodes communicate through TCP + SSL



Certificates are given in the genesis block


# The Red Belly Blockchain

The genesis block also contains a list of  $n$  participants



# The Red Belly Blockchain


External nodes access the blockchain through these participants



A tx is committed if  $t+1$  participants say so.

# The Red Belly Blockchain


The  $n$  nodes running the consensus...



are regularly changed:  $n, n', n'' \dots$  but  $t' < n'/3, t'' < n''/3 \dots$

# The Red Belly Blockchain

And encoded in the next block



if the correct  $n$  participants agree (consensus)



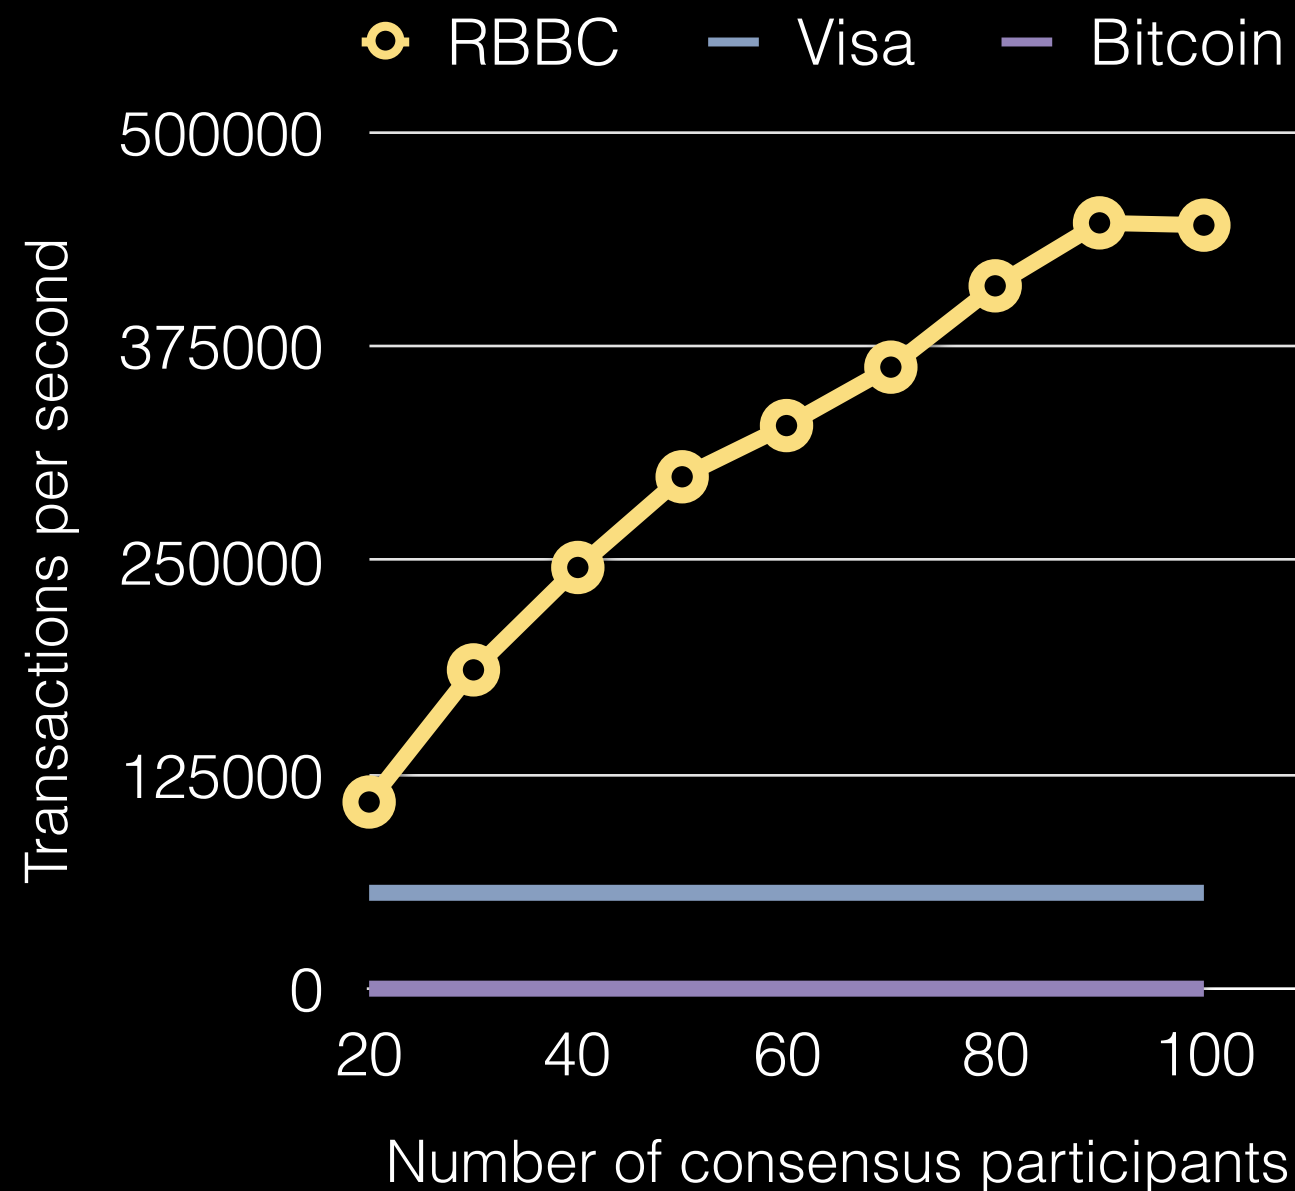
How does it perform?

# Experimental Settings

- Initiator sends a message to all to start
- Each node connects to each other through TLS
- 50 instances of consensus in which:
  - Each node proposes a block of 10K txs
  - Each tx is a 300-byte UTXO transaction
  - Each node spawns  $n$  instances of RBbcast and BBC
  - and quasi-validates each transaction (0.3msec)
  - Once  $n-t$  decided values 1

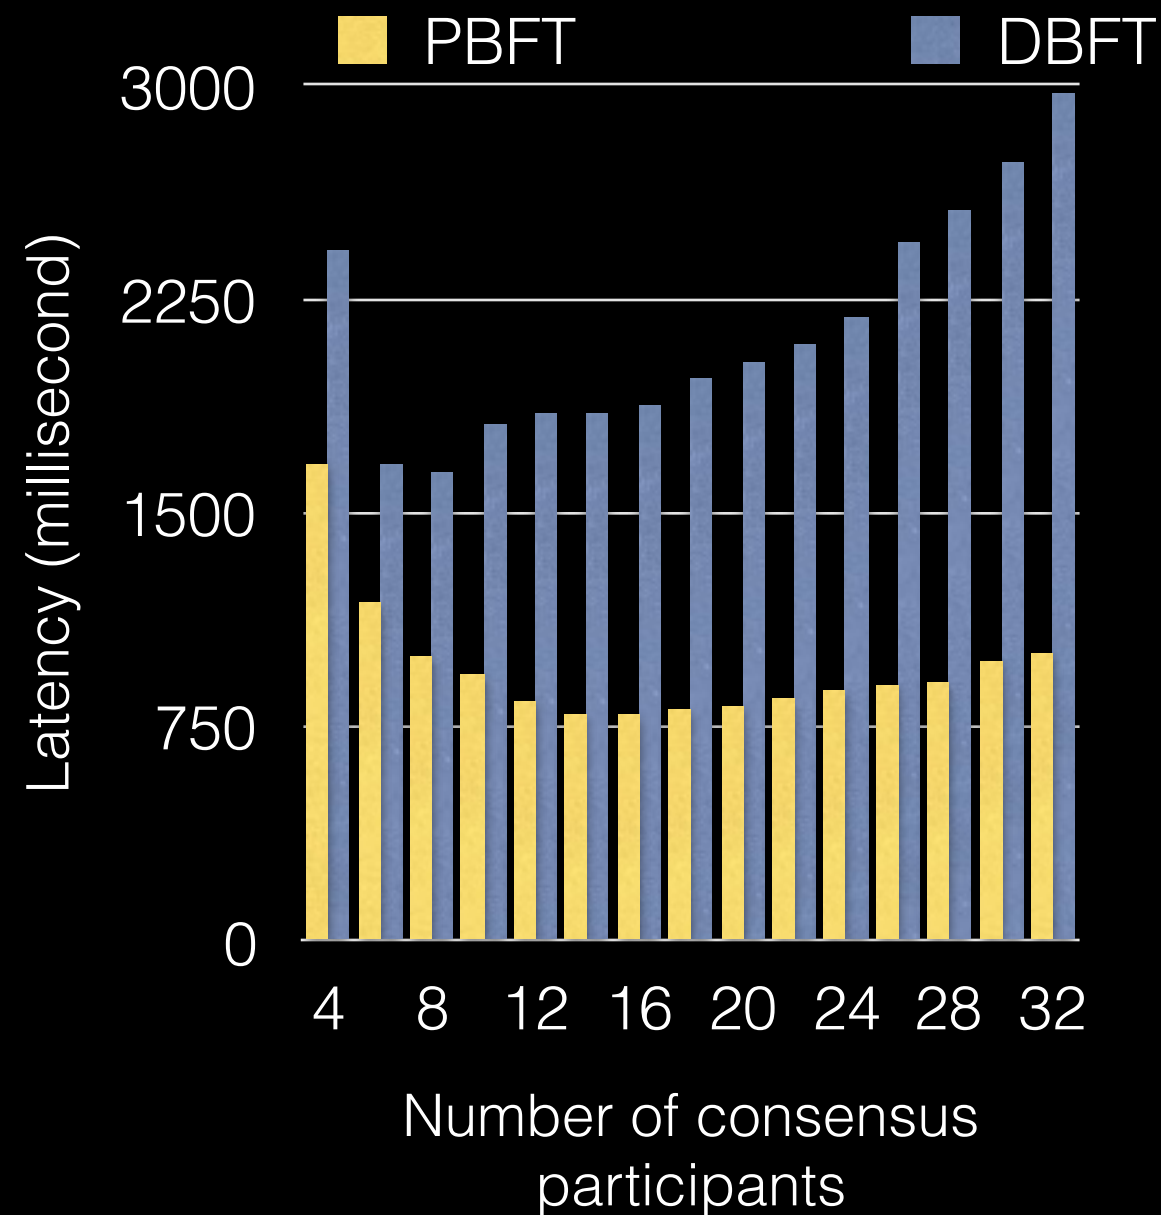
# Scalability

Amazon EC2, c4.4xlarge, 16vCPU, 30GiB mem, 2Gbps, t=6



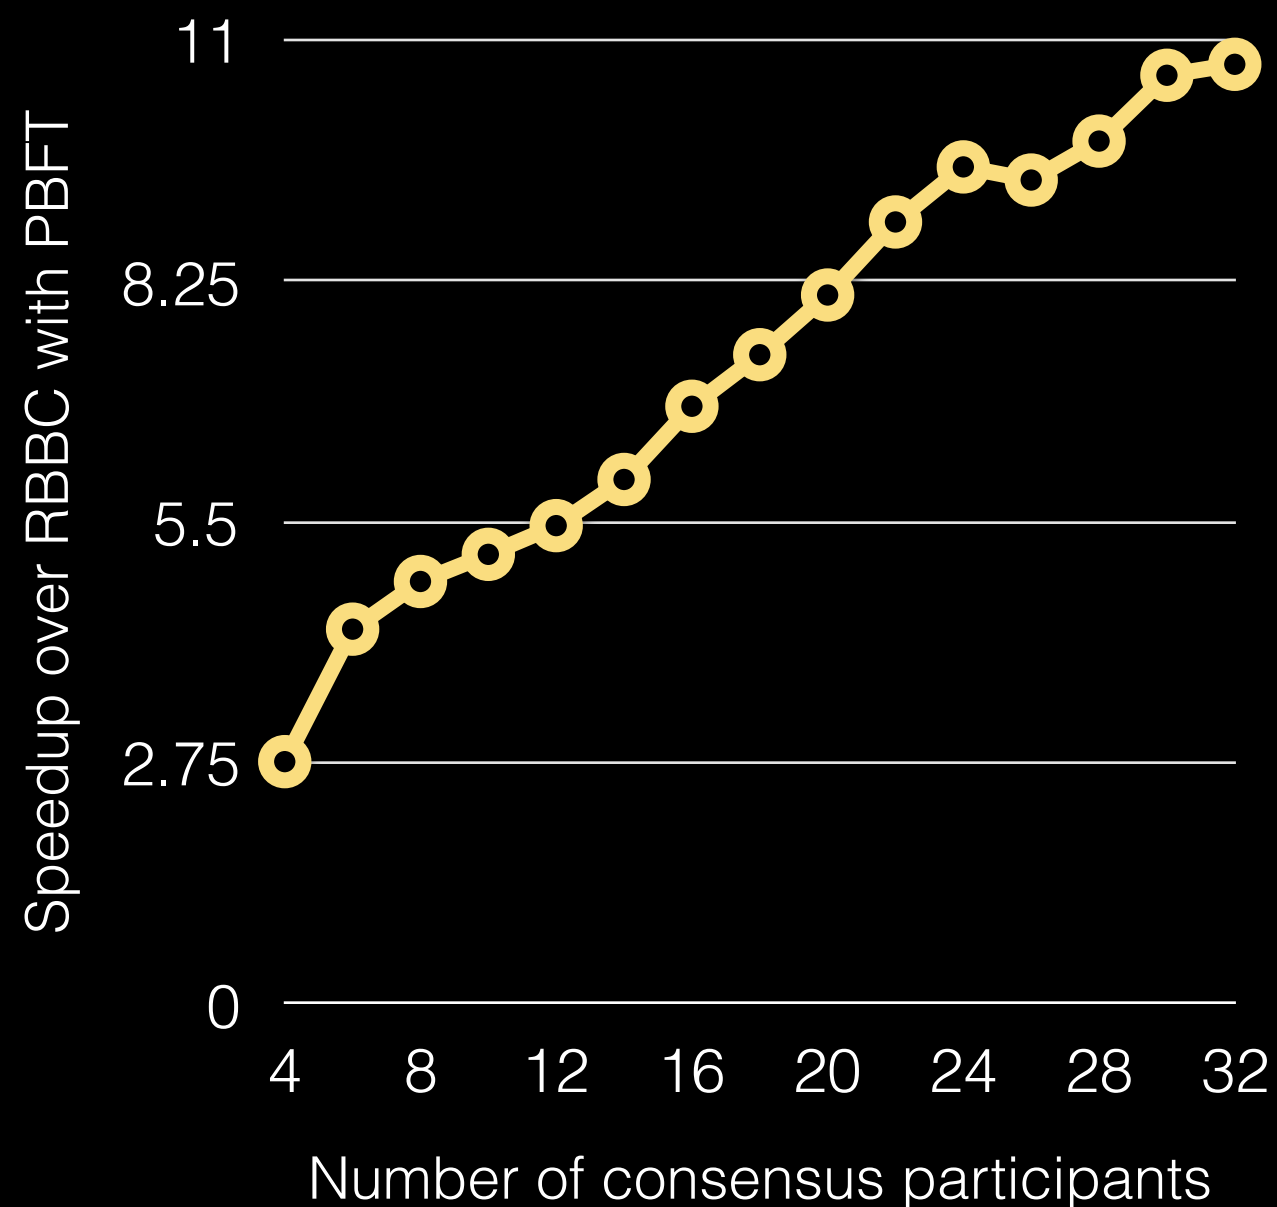
# Comparison with PBFT [CL02]

Intel E5530 2.4GHz 8CPU, 1Gbps, t=1




# Comparison with PBFT [CL02]

Intel E5530 2.4GHz 8CPU, 1Gbps,  $t=1$



# Number of Failures

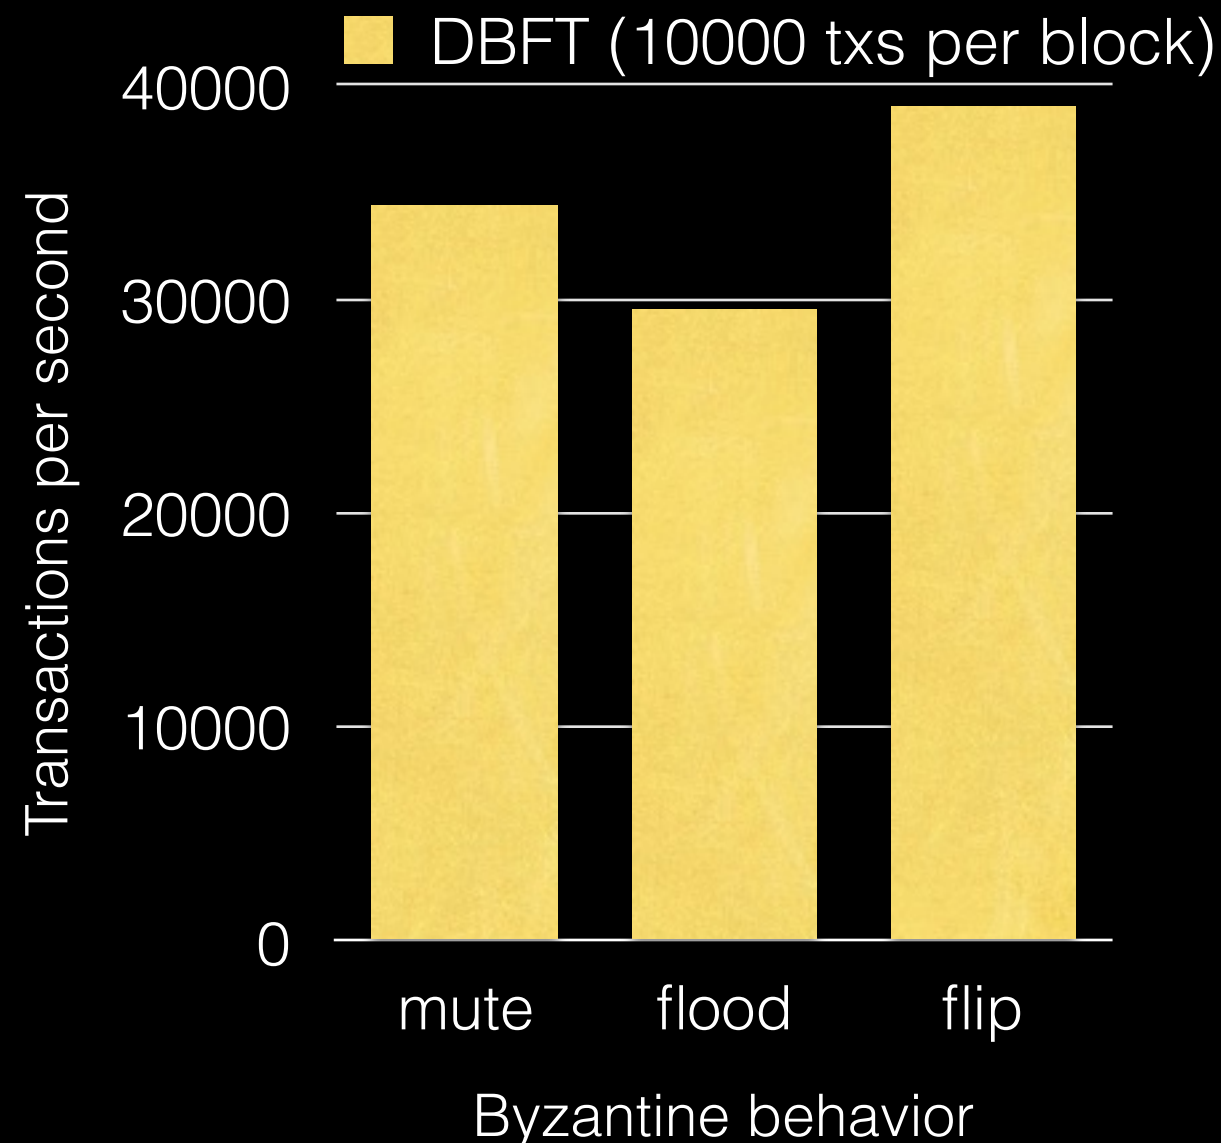
Intel E5530 2.4GHz 8CPU, 1Gbps, n=32





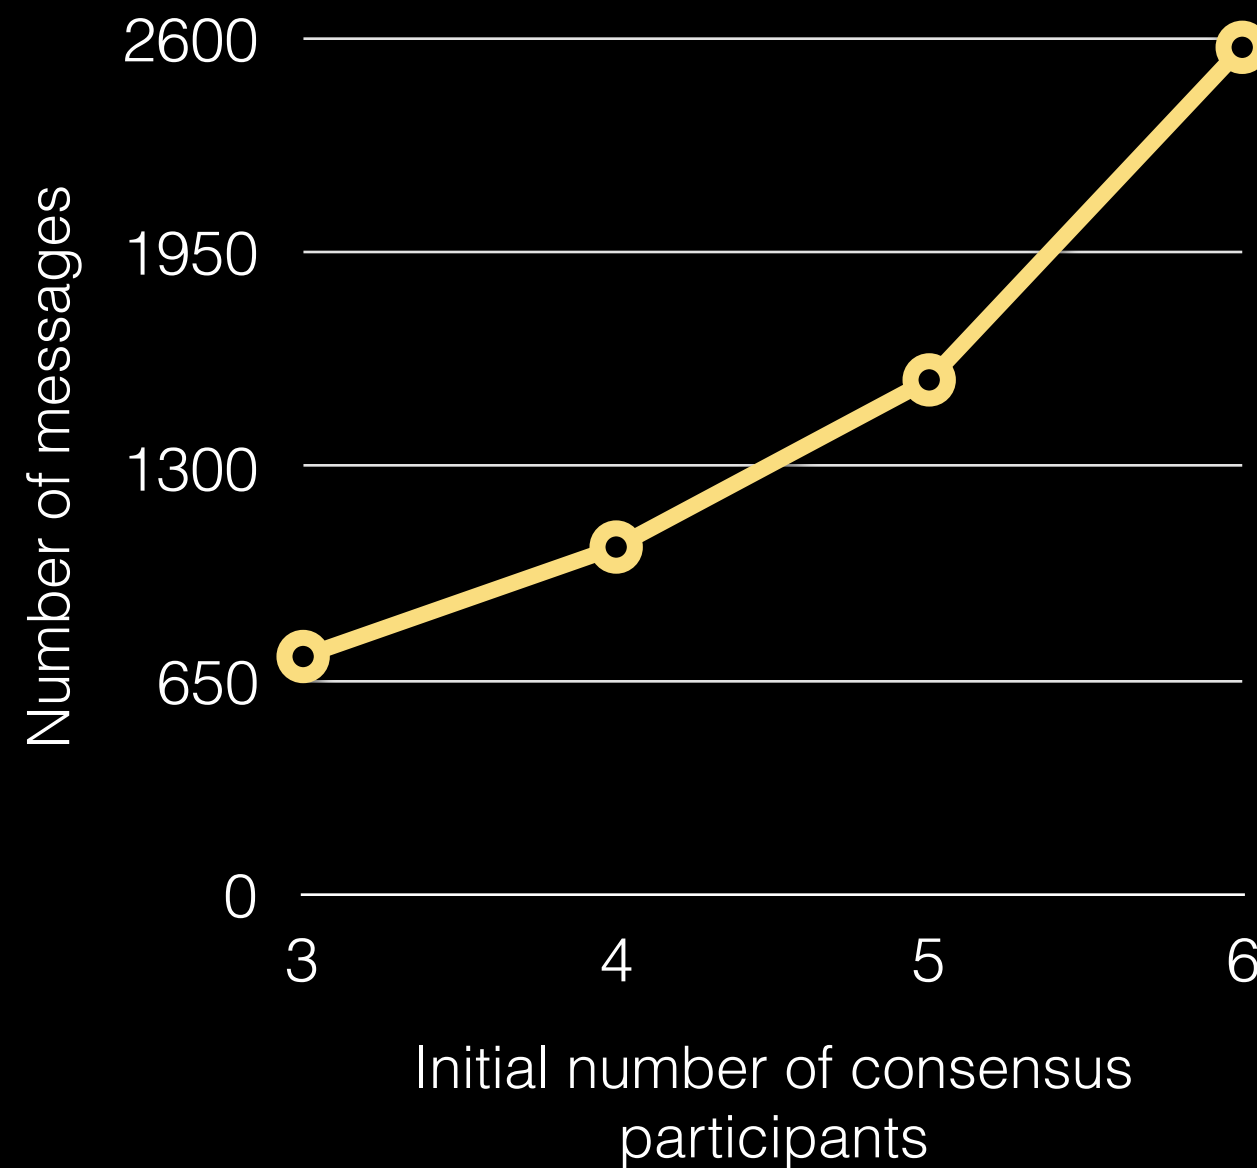
# Failures types

Intel E5530 2.4GHz 8CPU, 1Gbps, n=32, t=10



# Dynamism

Byzantine fault tolerant reconfiguration to 7 nodes



# Conclusion

- Blockchain systems provide **guarantees** not well-understood
  - This may lead to **dramatic** anomalies (and then double-spending attack)
- We propose the Red Belly Blockchain
  - Efficient: 400+ KTPS
  - Scale to 90 server nodes
  - Deterministic: with a new leaderless Byzantine consensus algorithm
  - Dynamic: A membership reconfiguration
- The next steps is to deploy it and build applications on top

# References

- [Nak'08] S. Nakamoto, "Bitcoin: a peer-to-peer electronic cash system," 2008. <http://www.bitcoin.org>
- [DSN'17] C. Natoli, V. Gramoli.
- [AlgoTel'17] T. Crain, V. Gramoli, M. Larrea, M. Raynal. Blockchain Consensus. May 2017.
- [CGLR'17] T. Crain, V. Gramoli, M. Larrea, M. Raynal. (Leader, randomization, signature)-tree Byzantine consensus for consortium blockchain. arXiv.
- [JACM'15] A. Mostefaoui, M. Raynal. Signature-Free Asynchronous Binary Byzantine Consensus with  $t < n/3$ ,  $O(n^2)$  Messages, and  $O(1)$  Expected Time. J. ACM 2015.
- [PODC'94] Asynchronous Secure Computations with Optimal Resilience. pp. 183–192. PODC 1994.
- [Bra'87] Bracha G., Asynchronous Byzantine agreement protocols. Information & Computation, 75(2):130–143, 1987.
- [CL'02] M. Castro and B. Liskov. Practical byzantine fault tolerance and proactive recovery, ACM Trans. Comput. Syst., vol. 20, no. 4, pp. 398– 461, 2002.



# See you in Sydney



# More information

<https://goo.gl/sV5GMC>